

# Comprehensive Literature Review on the Model Context Protocol (MCP)

## Section 1: Introduction to MCP (Synthesis)

### 1.1 The Agentic Paradigm Shift and the Integration Challenge

The evolution of Large Language Models (LLMs) represents a fundamental paradigm shift within artificial intelligence, moving models beyond passive text generation toward active, autonomous agency capable of executing real-world tasks. This agentic transformation necessitates robust and scalable mechanisms for external tool invocation.<sup>1</sup> Historically, the integration of LLMs with external capabilities was hampered by the **integration challenge**. In this scenario, LLM platforms required custom, hardcoded bindings for every one of available external tools or APIs. This resulted in distinct integration pathways, leading to excessive maintenance overhead, redundancy, and significant barriers to ecosystem scaling.<sup>2</sup>

The Model Context Protocol (MCP) was specifically developed to resolve this integration bottleneck. The protocol dictates a shift from framework-specific, application-centric tool bindings toward an interoperable ecosystem comprising composable and dynamically discoverable network services.<sup>2</sup> By standardizing the interface between the LLM and the external world, MCP eliminates the need for duplicate maintenance efforts and promotes a shared, scalable ecosystem for tool-augmented AI.<sup>2</sup>

### 1.2 Defining the Model Context Protocol (MCP): The "Universal Connector"

The Model Context Protocol (MCP), introduced by Anthropic in late 2024, is an open-source,

schema-driven standard that provides a consistent mechanism for AI systems to access external data, APIs, and tools outside their core model boundaries.<sup>1</sup> Often likened to a "universal connector" for AI, MCP is designed for real-time decision-making and is foundational to building scalable, intelligent agentic workflows.<sup>3</sup>

The architectural implementation of MCP is significant because it enforces a strict separation between the core LLM reasoning domain (the client) and the tool's execution environment (the server).<sup>4</sup> This decoupling is vital for ensuring architectural resilience and modularity. New tools can be added or existing tools updated without requiring corresponding changes to the client (agent) or host code, thereby allowing LLMs to flexibly extend their functionality by connecting to new servers on demand.<sup>2</sup> This enforced separation ensures that the complex logic of the agent remains stable even as the underlying external tool inventory changes and grows.

### 1.3 Structure and Scope of the Comprehensive Review

This review synthesizes findings from recent scholarly articles published post-2024, focusing exclusively on academic papers that detail MCP's architecture, adoption dynamics, empirical performance, and associated research challenges. The subsequent sections will detail the core architectural components, analyze the breakthroughs in implementation automation, review application areas in general and specialized domains, and conclude with a discussion of critical unresolved issues related to agentic reliability, security, and governance.

## Section 2: Key Academic Papers (Abstract List)

The following abstracts represent the core literature driving current research and understanding of the Model Context Protocol, focusing on its development, implementation, and empirical evaluation.

1. Abstract <sup>1</sup>: Large Language Models (LLMs) are increasingly evolving from passive text generators into active agents that can...[source](#)
2. Abstract <sup>5</sup>: Tool calling has emerged as a critical capability for AI agents to interact with the real world and solve complex...[source](#)
3. Abstract (Summary of Findings) <sup>2</sup>: Future research directions for the Model Context Protocol (MCP) are focused on strengthening its standardization, trust boundaries, and sustainable growth. Key challenges identified include security, scalability, and

governance issues. The independence of MCP server management by individual developers means there is no central authority to audit security baselines or enforce uniform compliance, leading to inconsistent patching and configuration drift. Future research should focus on strengthening trust boundaries, mandatory configuration validation, automated version checks, integrity auditing, and continuous community participation to ensure ecosystem-wide resilience. The expansion of server interactions across multiple lifecycle stages amplifies ethical and safety concerns, requiring solutions to ensure fairness in tool selection, defense against dataset leakage, and maintaining accountability in automated decision workflows.

4. Abstract <sup>6</sup>: LLMs' capabilities are enhanced by using function calls to integrate various data sources or API results into the...[source](#) the results.
5. Abstract (Economic Research Application) <sup>4</sup>: The objective of this paper is to demystify AI agents - autonomous LLM-based systems that plan, use tools,...[source](#) analysis.

## Section 3: Thematic Summary of Research (Synthesis)

### 3.1 Core Definition and Architecture

#### 3.1.1 Architectural Foundations: Client-Server Model and Protocol Design

The Model Context Protocol establishes a foundational client-server architecture. MCP clients are typically AI agents or applications that connect to servers to access capabilities.<sup>4</sup> These clients maintain research context across extended projects and execute the planning necessary for tool use.<sup>4</sup> Conversely, MCP servers are responsible for offering resources, hosting tools, and executing the actual API interactions with external data sources.<sup>4</sup>

The protocol's design is based on the JSON-RPC 2.0 standard, which ensures implementation consistency and reusability.<sup>3</sup> This choice facilitates a safety-first architecture, incorporating features essential for secure, autonomous operation. These features include strong typing, clearly defined request/response lifecycles, integrated permission layers, and mechanisms for flowing client-server alarms designed to prevent insecure or unintended auto-invocation.<sup>3</sup> This structural emphasis on defined interaction and typing is crucial for the deployment of

LLMs in production environments where reliability and security are non-negotiable.

### 3.1.2 Key Components and Schema Reliance

The MCP fundamentally relies on a schema-driven standard to dictate how external tools are described for dynamic discovery and invocation by the LLM.<sup>1</sup> The primary technical component that defines the external functionality is the schema contract. Academic literature confirms that the most effective method for generating these schemas and implementing MCP servers utilizes pre-existing industry standards, specifically OpenAPI 2.0/3.0 specifications.<sup>1</sup>

The LLM client requires a comprehensive description of the tool's parameters, inputs, and expected outputs to effectively integrate it into its reasoning chain. The MCP server registers these definitions, allowing the LLM to access capabilities such as file systems, web crawlers, or financial data.<sup>6</sup> This schema reliance provides the necessary structure for the LLM to transition from merely generating text to generating valid function calls that operate against real-world systems.

Table 3.1: MCP Architectural Components and Functions

Component Role	Core Function	Underlying Standard/Protocol	Key Feature/Constraint
MCP Client (Agent)	Discovers, invokes, and integrates tool outputs into the LLM context.	JSON-RPC 2.0	Constrained by context window limits; must handle tool enumeration token length. <sup>6</sup>
MCP Server (Tool Host)	Exposes external capabilities (APIs, databases); manages execution and authentication.	Schema-driven (OpenAPI)	Requires high-quality source specifications; initially bottlenecked by manual scaffolding. <sup>1</sup>

Protocol Design	Provides a unified, standardized interface for tool definition and interaction.	Open Standard, JSON-RPC 2.0	Enables Modularity, Safety-first features (permissions), and Scalable Optimization (caching, batching). <sup>3</sup>
-----------------	---	-----------------------------	--

### 3.2 Implementation, Scalability, and Adoption Dynamics

#### 3.2.1 Quantifying the Manual Server Development Bottleneck

Despite the protocol's goal of accelerating integration, initial adoption research revealed significant friction in implementing MCP servers. An analysis of over 22,000 MCP-tagged GitHub repositories created within six months of the protocol's release showed that fewer than 5% included functional server implementations.<sup>1</sup> These existing servers were typically small, single-maintainer projects characterized by repetitive manual effort, including writing glue code, handling authentication, and configuring schemas by hand.<sup>1</sup> This required developers to replicate much of the integration work MCP was intended to eliminate, demonstrating a critical initial threat to achieving the necessary network effects for widespread adoption.

#### 3.2.2 Automation Breakthrough: AutoMCP and the Role of OpenAPI

To overcome the costly bottleneck of manual server development, the concept of automated compiler solutions emerged. The AutoMCP compiler successfully demonstrated the ability to parse REST API definitions from OpenAPI 2.0/3.0 specifications and generate complete MCP server implementations, including necessary schema registration and authentication handling.<sup>1</sup>

Empirical evaluation of AutoMCP on 50 real-world APIs, encompassing 5,066 endpoints across more than 10 domains, yielded substantial success. From a stratified sample of 1,023

tool calls, 76.5% succeeded out-of-the-box.<sup>1</sup> Following minor fixes, necessitated by inconsistencies in the source OpenAPI contracts and averaging just 19 lines of spec changes per API, the success rate rose to an impressive 99.9%.<sup>1</sup>

### **3.2.3 Analysis of Specification Quality as the New Adoption Bottleneck**

The results from the automation analysis carry a profound implication for the technical focus of MCP adoption. The near-complete success of automated server generation indicates that the complexity inherent in LLM integration has fundamentally shifted: the challenge is no longer protocol design or code generation. Instead, the primary bottleneck for ecosystem expansion is now the quality and consistency of the underlying OpenAPI contracts used to define the servers.<sup>1</sup> Since failures almost exclusively resulted from specification omissions or inaccuracies, organizations wishing to capitalize on the MCP ecosystem must prioritize API governance and documentation accuracy. The protocol's success has inadvertently created a pressure point on legacy IT practices surrounding specification maintenance, demanding greater rigor in how API contracts are authored and maintained.

## **3.3 Primary Application Areas and Real-World Use Cases**

### **3.3.1 General Agentic Workflows and Ecosystem Growth**

MCP has rapidly cemented its status as a foundational architecture for AI-native applications. This is demonstrated by the thousands of independently developed MCP servers that expose model-accessible interfaces to ubiquitous enterprise and consumer services such as GitHub and Slack.<sup>2</sup> Benchmarks like MCPToolBench++ draw on a marketplace of over 4,000 MCP servers spanning more than 40 distinct categories, validating the protocol's wide-ranging applicability across tasks including data analysis, file operations, financial calculations, and general compute.<sup>6</sup> Client applications can flexibly extend their functionality simply by connecting to new servers as required.<sup>2</sup>

### **3.3.2 Specialized Domain Application: Economic and Institutional Research**

One high-value application area is specialized research within high-trust environments, such as economics or institutional data analysis. MCP enables AI agents to connect to and maintain persistent connections with institutional databases (e.g., those held by central banks or private research groups).<sup>4</sup>

This capability is transformative because it enables complex research workflows, such as autonomously conducting literature reviews, writing and debugging econometric code, and fetching and analyzing proprietary economic data.<sup>4</sup> Critically, MCP serves as a powerful abstraction layer: it handles the complex integration details, allowing the end-user (the economist or domain expert) to command the agent using natural language, or "vibe coding," without needing specialized knowledge of the institutional API details.<sup>4</sup> By transferring integration complexity to the standardized server design, the protocol democratizes the use of sophisticated, high-value proprietary data sources for autonomous research workflows.

## 3.4 Research Focus: Empirical Benchmarking and Performance Analysis

### 3.4.1 Introduction to State-of-the-Art Benchmarks

Two key benchmarks have emerged to rigorously evaluate the performance of MCP-enabled agents. **LiveMCP-101** provides 101 carefully curated real-world queries demanding the coordinated use of multiple diverse MCP tools (e.g., web search, file operations, reasoning). Its novel evaluation approach leverages ground-truth execution plans rather than relying solely on raw API outputs, better reflecting realistic, dynamic execution scenarios.<sup>5</sup>

**MCPToolBench++** addresses the challenges posed by the diverse response formats from different MCP server executions and the inherent variability in tool success rates in real-world settings.<sup>6</sup> This benchmark is built upon the expansive MCP server marketplace, offering a multi-domain framework for evaluating both single-step and multi-step tool calls.<sup>6</sup>

### 3.4.2 Performance Findings: The Tool Orchestration Deficit

Despite the significant architectural standardization provided by MCP, performance evaluations confirm a fundamental limitation in agentic reasoning. Experiments using the LiveMCP-101 benchmark demonstrate that even the most advanced, frontier LLMs achieve a success rate below 60% on complex, multi-step tool orchestration tasks.<sup>5</sup> This low rate of reliability confirms that while MCP successfully standardizes *how* tools are accessed, this standardization is necessary but insufficient to guarantee *reliable execution*; the core limitation resides in the LLM's long-range planning, coordination, and reasoning capabilities required to navigate dynamic environments.

### 3.4.3 Qualitative Analysis of Failure Modes and Resource Constraints

Detailed error analysis has identified distinct failure modes in LLMs when utilizing MCP tools. A significant category is Tool Orchestration Errors, including the critical failure mode known as "overconfident self-solving".<sup>5</sup> In this scenario, the agent recognizes the requirement for an external tool but chooses to rely instead on its own internal knowledge or reasoning, bypassing the externally verified and grounded MCP tool. This often results in generic or hallucinated answers and premature termination of the task.<sup>5</sup> Other orchestration errors include outright ignoring explicitly stated requirements or failing to select the relevant tool.<sup>5</sup> Implementation failures often manifest as Parameter Errors, where the agent inaccurately formats or omits required input parameters necessary for the successful execution of the MCP tool call.<sup>5</sup>

Beyond errors in execution logic, the research highlights a critical constraint tied to the LLM's context window. The textual descriptions of tools and their parameters necessary for MCP server invocation consume significant token length.<sup>6</sup> As the ecosystem scales (now over 4,000 servers), listing a comprehensive inventory of available schemas directly reduces the available token space needed for complex planning, reasoning, and processing multi-step outputs.<sup>5</sup> This establishes a necessary resource allocation trade-off between inventory size and reasoning depth, compromising performance in complex tasks.<sup>5</sup>

Table 3.2: Observed Failure Modes in MCP-Enabled Agent Execution (LiveMCP-101)

Error Category	Example Failure Mode	Detailed Description	Source
Tool Orchestration	Success Rate Below 60%	Empirical evidence demonstrating LLM	<sup>5</sup>



		failure in coordinating complex, multi-step actions in realistic environments.	
Tool Orchestration	Overconfident Self-Solving	Agent recognizes requirement but chooses internal reasoning/knowledge, resulting in hallucinated outputs and bypassing the grounded MCP tool.	5
Tool Orchestration	Ignoring Requirement	Agent misses an explicitly stated requirement and does not select the relevant tool, leading to early termination or a generic final answer.	5
Implementation	Parameter Errors	Agent inaccurately formats or omits required input parameters necessary for a successful MCP tool call execution.	5
Scalability/Context	Token Inefficiencies/Limits	Tool enumeration (schemas) depletes the context window, forcing a resource allocation trade-off that	5

		compromises complex reasoning.	
--	--	--------------------------------	--

## Section 4: Conclusion and Research Gaps (Synthesis)

### 4.1 Summary of the Current State of MCP Research

The Model Context Protocol has fulfilled its primary objective of standardizing the architectural foundation for LLM-tool interaction, effectively solving the previous integration problem.<sup>1</sup> Research demonstrates that the protocol's reliance on OpenAPI contracts, coupled with automated server generation, has dramatically reduced the technical hurdle for developers to expose external capabilities to AI agents.<sup>1</sup> The ecosystem has experienced rapid growth, validating MCP as a key advancement in LLM architecture and enabling sophisticated applications in high-stakes domains like economic research.<sup>3</sup>

However, the standardization achieved by MCP has exposed two resulting critical areas requiring immediate scholarly attention: fundamental limitations in agentic reliability and profound challenges in ecosystem governance.<sup>2</sup> The community has effectively transitioned its focus from solving tool *integration* problems to analyzing the resulting problems of autonomous *execution* and *sustainability*.

### 4.2 Major Unresolved Challenges and Future Directions

#### 4.2.1 Security Vulnerabilities and the Need for Robust Trust Boundaries

The success and decentralized nature of MCP introduce inherent systemic risks related to security. The independence of MCP server management means there is no central authority to audit security baselines or enforce uniform compliance.<sup>2</sup> This decentralization fosters heterogeneity across deployment and maintenance, leading to inconsistent patching, irregular

security practices, and configurations that drift over time, increasing the likelihood of vulnerabilities.<sup>2</sup>

Future research must prioritize establishing and strengthening trust boundaries across the decentralized ecosystem.<sup>2</sup> This requires the implementation of technical governance solutions, such as integrating mandatory configuration validation, automated version checks, and integrity auditing into MCP registries and collections. Without these mechanisms, the sustainable growth of the ecosystem is threatened by fragmentation and the potential for widespread security compromises due to non-uniform baseline compliance.<sup>2</sup>

#### **4.2.2 Scalability Issues, Fragmentation, and Governance Requirements**

The hard constraint of the LLM's context window limits the number of available tools that an agent can effectively consider in a single run, directly curtailing the scalability of complex, tool-heavy workflows.<sup>6</sup> Future research must prioritize dynamic, contextual tool discovery mechanisms and innovative schema compression techniques to mitigate the trade-off between maintaining a comprehensive tool inventory and reserving sufficient token space for complex reasoning and planning.<sup>6</sup>

Furthermore, as MCP facilitates the expansion of agent interaction into high-stakes institutional domains (e.g., accessing financial models or proprietary datasets), the existing challenges related to low reliability (sub-60% success rate on multi-step tasks) are amplified.<sup>4</sup> If an agent makes an error in orchestration, or if it accesses a critical resource through a poorly secured server, the resulting failure involves significant ethical, safety, and legal ramifications. Therefore, future research must address governance to ensure responsible development. Shared priorities include ensuring fairness in the agent's selection of tools, defending against potential dataset leakage, and, most critically, maintaining clear accountability chains for automated decision workflows. Establishing governance protocols that match the technical capabilities of MCP is paramount to ensuring the protocol's long-term sustainable growth and responsible application.<sup>2</sup>

#### **Alıntılanan çalışmalar**

1. Making REST APIs Agent-Ready: From OpenAPI to MCP ... - arXiv, erişim tarihi Ekim 13, 2025, <https://arxiv.org/abs/2507.16044>
2. Model Context Protocol (MCP): Landscape, Security Threats ... - arXiv, erişim tarihi Ekim 13, 2025, <https://arxiv.org/pdf/2503.23278>
3. What Is Model Context Protocol (MCP) | How it Works - Kodexo Labs, erişim tarihi Ekim 13, 2025, <https://kodexolabs.com/what-is-model-context-protocol-mcp/>
4. NBER WORKING PAPER SERIES AI AGENTS FOR ECONOMIC ..., erişim tarihi Ekim

- 13, 2025, [https://www.nber.org/system/files/working\\_papers/w34202/w34202.pdf](https://www.nber.org/system/files/working_papers/w34202/w34202.pdf)
5. LiveMCP-101: Stress Testing and Diagnosing MCP-enabled ... - arXiv, erişim tarihi Ekim 13, 2025, <https://arxiv.org/abs/2508.15760>
  6. MCPToolBench++: A Large Scale AI Agent Model Context ... - arXiv, erişim tarihi Ekim 13, 2025, <https://arxiv.org/abs/2508.07575>